

# JCoz: A Causal Java Profiler

By David Vernet and Matt Perron



# Example program with 32 threads

```
public static void main(String[] args) throws InterruptedException {
    for (int i = 0; i < numThreads; i++)
        threads.add(new ParallelWorker());

    while (true) {
        doParallel();
        doSerial();
        System.out.println("Iteration done");
    }
}



public static void doSerial() throws InterruptedException {
    long sum = 0;
    for (long i = 0; i < LOOP_ITERS; i++)
        sum += (System.nanoTime() % 9999);
}
```

```
public static void doParallel() throws InterruptedException {
    executor.invokeAll(threads);
}

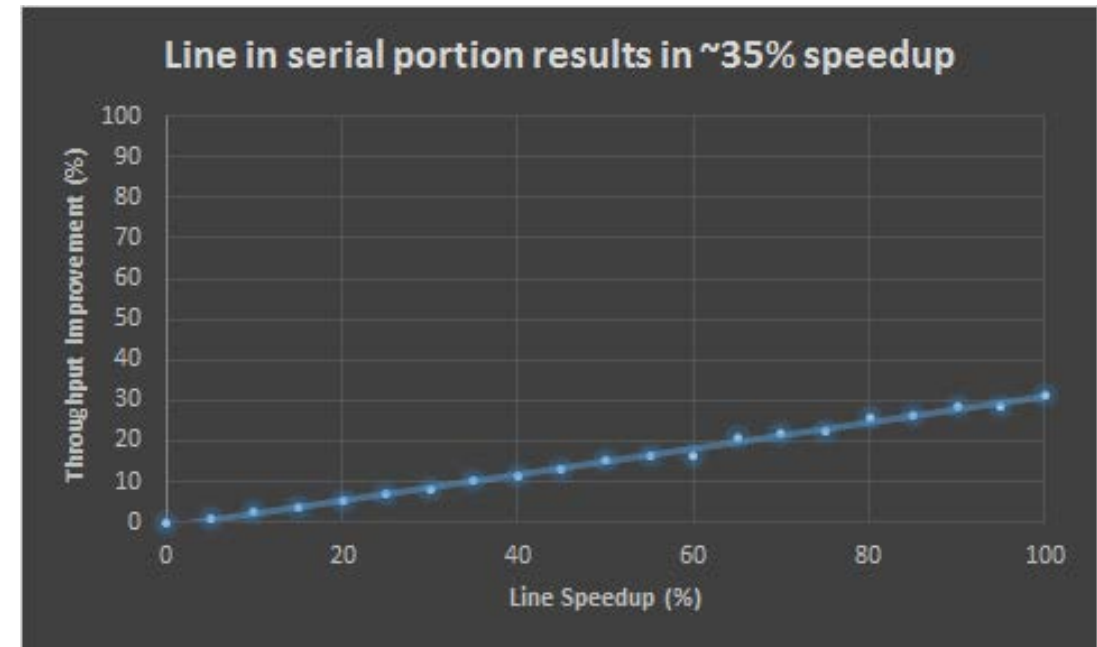
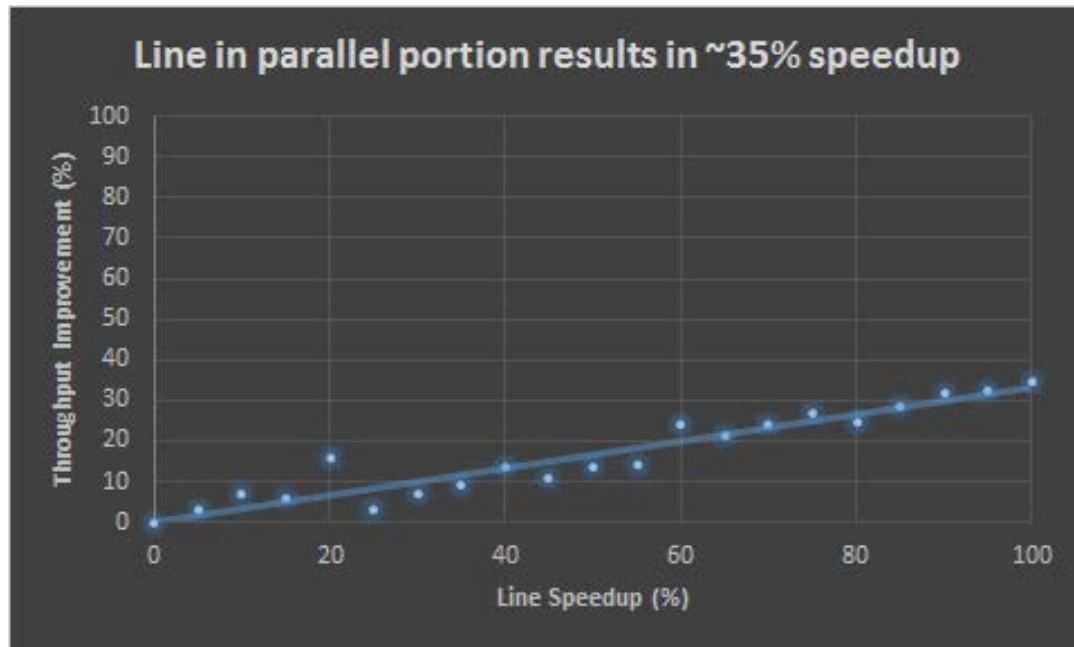
static class ParallelWorker implements Callable<Void> {
    public Void call() {
        long sum = 0;
        for (long i = 0; i < LOOP_ITERS; i++)
            sum += (System.nanoTime() % 9999);

        return null;
    }
}
```

# Standard sampling profiler output gives misleading results

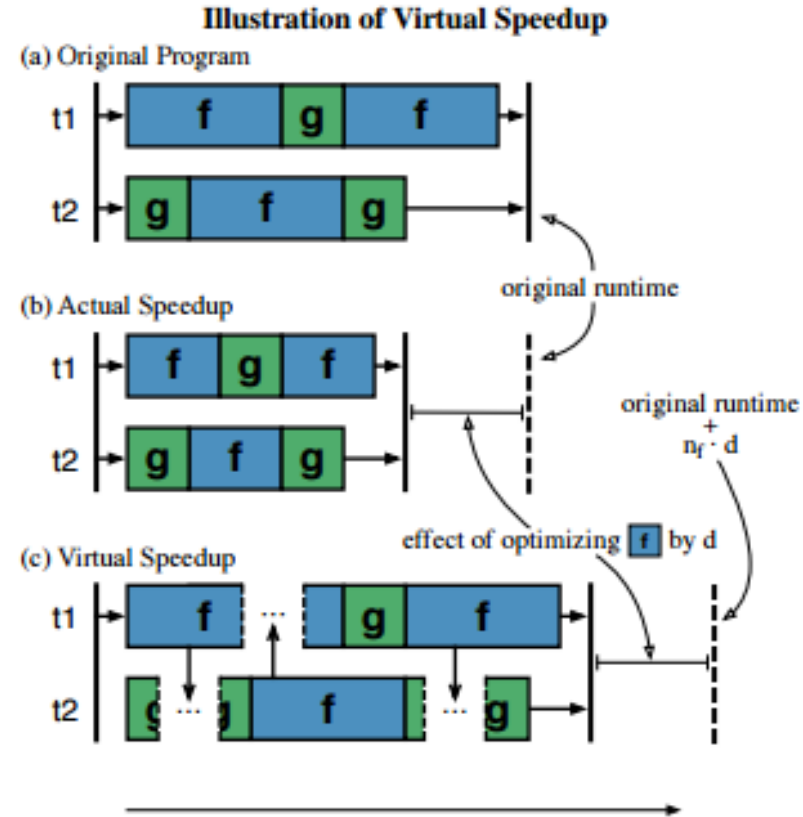
Hot Spots - Method	Self Time [%]	Self Time	Self Time (CPU) ▼	Total Time	Total Time (CPU)
test.TestThreadSerial\$ThreadTest. <b>call</b> ()		1,532,364 ms (97.5%)	1,532,364 ms	1,532,364 ms	1,532,364 ms
test.TestThreadSerial. <b>doSerial</b> ()		39,445 ms (2.5%)	39,445 ms	39,445 ms	39,445 ms
test.TestThreadSerial. <b>doParallel</b> ()		57,550 ms (0%)	99.2 ms	57,550 ms	99.2 ms
test.TestThreadSerial. <b>main</b> ()		0.000 ms (0%)	0.000 ms	96,996 ms	39,544 ms

# JCoz gives accurate profile results



How does causal  
profiling work?

# Virtual Speedup



**Figure 3:** An illustration of virtual speedup: (a) shows the original execution of two threads running functions  $f$  and  $g$ ; (b) shows the effect of a *actually* speeding up  $f$  by 40%; (c) shows the effect of *virtually* speeding up  $f$  by 40%. Each time  $f$  runs in one thread, all other threads pause for 40% of  $f$ 's original execution time (shown as ellipsis). The difference between the runtime in (c) and the original runtime plus  $n_f \cdot d$ —the number of times  $f$  ran times the delay size—is the same as the effect of actually optimizing  $f$ .

Photo credit: Charlie Curtsinger and Emery Berger

C. Curtsinger, E. Berger. COZ: Finding Code that Counts with Causal Profiling. *SOSP '15 ACM SIGOPS*

# Virtual Speedup implementation

- Throughout runtime of program, run “experiments”
  - Randomly choose a recently executed line
  - Randomly choose a speedup between 0-100%
  - Every time a thread hits that line in a sample, freeze other threads for the speedup threshold chosen above
  - Measure how much throughput is achieved for the given experiment line and speedup
- At the end of the program (when enough results have been collected), experiments that indicate biggest throughput increases relative to the experiment runtime are highlighted as lines for optimization


# Library optimizations with JCoz



# Optimized Java H2 Database engine

- Using JCoz, we were able to optimize the widely used, mature, Java H2 Database Engine: <http://www.h2database.com/html/main.html>
- Found and implemented optimization within 2 hours of running profiler
- No previous knowledge of codebase
- Measured optimization of TPCC using the Dacapo Benchmark Suite
- Used in many projects, including Apache Cayenne, Apache Jackrabbit, Jboss Jopr, and the NIH

# JCoz profile output for a line in H2 codebase



# Profiled line was excessive sleep on transaction failure


```
Database database = session.getDatabase();
int sleep = 1 + MathUtils.randomInt(10); ←
while (true) {
    try {
        if (database.isMultiThreaded()) {
            Thread.sleep(sleep); ←
        } else {
            database.wait(sleep);
        }
    } catch (InterruptedException e1) {
        // ignore
    }
    long slept = System.nanoTime() / 1000000 - now;
    if (slept >= sleep) {
        break;
    }
}
```

# Simple fix – sleep for less time

```
Database database = session.getDatabase();

// Sleep for 1-10 microseconds instead of 1-10 milliseconds
int sleep = 1000 * (1 + MathUtils.randomInt(10));
while (true) {
    try {
        if (database.isMultiThreaded()) {
            Thread.sleep(0, sleep);
        } else {
            database.wait(0, sleep);
        }
    } catch (InterruptedException e1) {
        // ignore
    }
    long slept = System.nanoTime() - now;
    if (slept >= sleep) {
        break;
    }
}
```

# Optimization successful – 20% speedup



# JCoz produced expected results

- No prior knowledge of codebase needed
- Pointed to the exact line of bottleneck
- Optimization was not made obvious using other profilers

# Conclusion

# JCoz works

- JCoz is the first causal Java profiler
- Provides more valuable information than existing profilers for multithreaded Java programs
- Allows profiling of throughput rather than just runtime
- Low runtime overhead



# Future work

- Continue optimizing libraries
  - Another optimization opportunity discovered in the Universal Java Matrix Package (UJMP) (poor cache locality on dense matrix multiplication)
- Add more features for tuning profiler
  - Add latency profiling
- Explore ways to extend causal profiling to multi-process and distributed applications

# Acknowledgements

- Special thanks to:
  - Charlie Curtsinger and Emery Berger for discovering causal profiling and making the original COZ program available for reference
  - David Capwell and Jeremy Manson for their code that we leveraged to use the infamous undocumented AsyncGetStackTraces JVMTI function
  - Professor Railing, for giving us guidance during this process

Questions?

# Appendix

# Another simple multithreaded example

```
public static void main(String[] args) throws InterruptedException {
    threads.add(new LongWorker());
    threads.add(new ShortWorker());




    while (true) {
        doParallel();
        System.out.println("Iteration done");
    }
}

public static void doParallel() throws InterruptedException {
    executor.invokeAll(threads);
}
```

```
static class LongWorker implements Callable<Void> {
    public Void call() {
        long sum = 0;
        for (long i = 0; i < 200000000L; i++)
            sum += (System.nanoTime() % 9999);
        System.out.println("Long done");
        return null;
    }
}

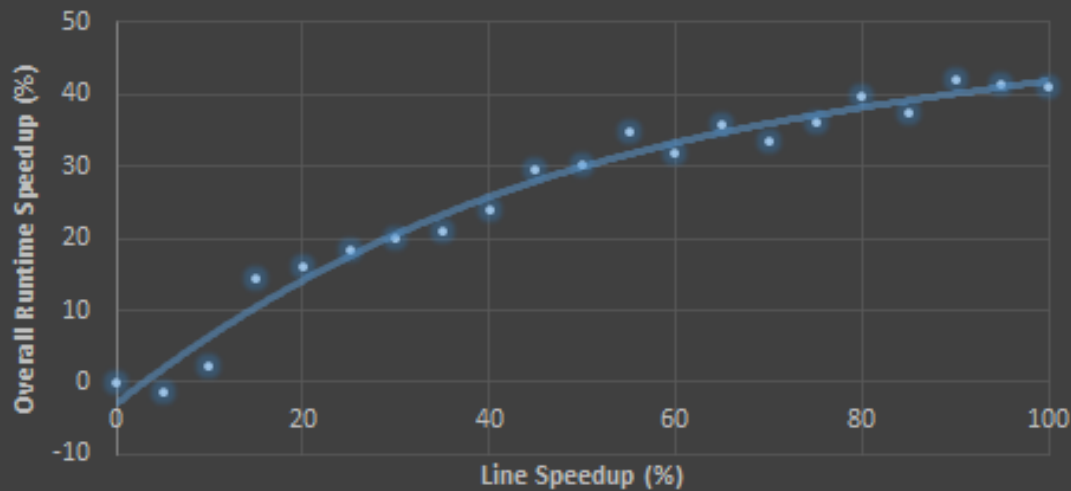
static class ShortWorker implements Callable<Void> {
    public Void call() {
        long sum = 0;
        for (long i = 0; i < 100000000L; i++)
            sum += (System.nanoTime() % 9999);
        System.out.println("Short done");
        return null;
    }
}
```

# Sampling profile again shows misleading results

Hot Spots - Method	Self Time [%] ▼	Self Time	Self Time (CPU)	Total Time	Total Time (CPU)
test.Test\$LongWorker. <b>call</b> ()		83,398 ms (63.7%)	83,398 ms	83,398 ms	83,398 ms
test.Test\$ShortWorker. <b>call</b> ()		42,672 ms (32.6%)	42,672 ms	42,672 ms	42,672 ms
test.Test. <b>doParallel</b> ()		4,695 ms (3.6%)	0.000 ms	4,695 ms	0.000 ms
test.Test. <b>main</b> ()		96.1 ms (0.1%)	96.1 ms	4,791 ms	96.1 ms

# JCoz gives accurate profile results

**Longer thread causal profile shows no benefit of speedup after 50%**



**Shorter Thread Causal Profile Shows No Speedup**

